

SANKET VISHWAKARMA

Data Engineer | Databricks • Snowflake • PySpark • AWS • CI/CD

vishwakarmasanket0@gmail.com | +91-7980956276 | Greater Noida, India

[linkedin.com/in/sanket-vishwakarma-902953109](https://www.linkedin.com/in/sanket-vishwakarma-902953109)



SUMMARY

Data Engineer with nearly 5 years building data platforms on AWS, Databricks, and Snowflake. I like owning a project end to end: scoping what it should do, planning the deployment, and staying with it through release. Day to day that means PySpark, SQL, DBT, and lakehouse work. Along the way I've cut manual effort by 65–75% through automation and CI/CD, and taken roughly \$100K a year out of our cloud bill.

TECHNICAL SKILLS

Data Engineering: PySpark, SQL, Python, ETL/ELT Design, Data Migration, Delta Lake, Batch & Stream Processing, Data Quality & Governance

Lakehouse & Platforms: Data Lakehouse Framework, Databricks (Lakeflow, SDF, Notebooks), Snowflake, AWS Lake Formation, DBT, Airflow, Terraform

Cloud & Infra (AWS): Elastic Beanstalk application infrastructure, Glue, Lambda, Step Functions, Appflow, RDS, SSL/SSO Integration

App Development: Streamlit, Node.js, Next.js, Python, PostgreSQL, Alembic (schema migrations)

CI/CD & Platform Admin: CI/CD design & optimization (auto-merge, branching, rule sets), Git, access management, cost & token optimization

AI / GenAI & Reporting: LLM-powered SQL generation, GenAI application design, ML, Statistical Analysis, Power BI

PROFESSIONAL EXPERIENCE

Data Engineer

Mars (following Kellanova acquisition)

Jan 2026 – Present

Greater Noida, India

- Ran the data migration from Kellanova's legacy platforms onto Mars infrastructure, working across Snowflake, Databricks, and AWS to close the transfer gaps that were blocking data flow between the two companies.
- Built the Elastic Beanstalk setup the team now uses to deploy their own Node.js, Next.js, and Python/Streamlit apps. 5+ applications are running on it in production.
- Shipped a Snowflake Streamlit app for data quality and governance. Users onboard a validation rule, an LLM writes the SQL for the check at run time, and the app reports back the results. Integrated with DBT, with dashboards and an approval workflow. Data quality improved 65%.
- Wrote Databricks Lakeflow jobs and models using SDF and notebooks to process data across the combined Mars and Kellanova platforms.
- Handle platform admin: SSL and SSO integrations, role-based access, CI/CD tuning, and cutting token usage to bring runtime cost down.
- Mentored interns through a 6-month program to completion. Also manage vendor partners on data engineering projects and sit on hiring panels.

Data Engineer

Sep 2023 – Dec 2025

Kellanova (formerly Kellogg's), Global Data Analytics Team

Mumbai, India

- Took an end-to-end orchestration POC through to deployment, building reusable CI/CD utilities and a structured, IAM-governed repo deployment model. Included a monitoring dashboard (Python, Step Functions, Lambda, APIs) that gave the org visibility into job status and cut manual monitoring by 70%.
- Owned the DevOps CI/CD pipelines, including auto-merge, branching strategy, and rule sets. Deployments got more reliable and we broke production less.
- Cut about \$100K a year off the AWS bill by auditing KMS key usage across objects, compacting small files, archiving cold data, and switching to bucket-level encryption.
- Brought Snowflake compute cost down after profiling how long warehouses were actually staying up, then consolidating jobs into a buffered scheduling window.
- Set up a CI/CD streaming process for DBT-to-Snowflake deployments that uses hooks to capture the full audit history for our dashboards.
- Wrote a data onboarding app on Alembic that handles config-metadata inserts and rollbacks in PostgreSQL, so nobody has to do schema changes by hand anymore.

Associate Data Engineer

Apr 2022 – Sep 2023

Edelweiss Tokio Life Insurance

Mumbai, India

- Built the agent-commission pipelines in AWS Glue and PySpark, pulling from multiple sources with delta loads and API integrations. Spark tuning got runtime down 40%.
- Led the POC and rollout of C-Data to replace a D365 ERP ingestor that was being decommissioned. Minimal downtime, nothing downstream broke.
- Saved \$2,000 a month by automating critical Glue jobs against the ERP sync database and clearing out orphan tables and unused jobs.

IT Programmer

Oct 2021 – Mar 2022

Cognizant

Kolkata, India

- Did statistical analysis and EDA across Macy's e-commerce data sources, and automated the web testing with Selenium and XPath, which halved manual testing time.

EDUCATION

B.Sc. Mathematics and Statistics — Jadavpur University

Jan 2015 – Nov 2020

Academic Excellence Award (Class XII)

CERTIFICATIONS

Databricks Fundamentals • Snowflake Data Application Builders • Apache Spark Big Data Engineering